

Kaggle Playground Series - Season 3, Episode 11

Jeremy Haakenson

2023-03-23

R Markdown

This file shows my approach to a Kaggle competition where the goal was to predict cost based on 16 variables.

Load packages.

```
library(splines)
library(gam)
```

```
## Loading required package: foreach
```

```
## Loaded gam 1.22-1
```

```
library(ranger)
library(gbm)
```

```
## Loaded gbm 2.1.8.1
```

```
library(e1071)
```

Load files.

```
train = read.csv('train.csv')
test = read.csv('test.csv')
```

Examine data.

```
summary(train)
```

```
##           id           store_sales.in.millions. unit_sales.in.millions.
##  Min.      :    0      Min.      : 0.510           Min.      :1.000
##  1st Qu.: 90084      1st Qu.: 3.720           1st Qu.:3.000
##  Median :180168      Median : 5.780           Median :3.000
##  Mean   :180168      Mean   : 6.337           Mean   :3.044
##  3rd Qu.:270251      3rd Qu.: 8.400           3rd Qu.:4.000
##  Max.    :360335      Max.    :22.920           Max.    :6.000
##  total_children num_children_at_home avg_cars_at.home.approx..1
##  Min.      :0.000      Min.      :0.0000           Min.      :0.000
##  1st Qu.:1.000      1st Qu.:0.0000           1st Qu.:1.000
```

```
## Median :2.000 Median :0.0000 Median :2.000
## Mean :2.456 Mean :0.6894 Mean :2.204
## 3rd Qu.:4.000 3rd Qu.:1.0000 3rd Qu.:3.000
## Max. :5.000 Max. :5.0000 Max. :4.000
## gross_weight recyclable_package low_fat units_per_case
## Min. : 6.00 Min. :0.0000 Min. :0.0000 Min. : 1.00
## 1st Qu.: 9.71 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:10.00
## Median :13.60 Median :1.0000 Median :0.0000 Median :20.00
## Mean :13.82 Mean :0.5681 Mean :0.3278 Mean :18.97
## 3rd Qu.:17.70 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:28.00
## Max. :21.90 Max. :1.0000 Max. :1.0000 Max. :36.00
## store_sqft coffee_bar video_store salad_bar
## Min. :20319 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:23593 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :27694 Median :1.0000 Median :0.0000 Median :1.0000
## Mean :28180 Mean :0.5648 Mean :0.2774 Mean :0.5048
## 3rd Qu.:33858 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:1.0000
## Max. :39696 Max. :1.0000 Max. :1.0000 Max. :1.0000
## prepared_food florist cost
## Min. :0.0000 Min. :0.0000 Min. : 50.79
## 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.: 70.32
## Median :1.0000 Median :1.0000 Median : 98.81
## Mean :0.5048 Mean :0.5032 Mean : 99.61
## 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:126.62
## Max. :1.0000 Max. :1.0000 Max. :149.75
```

```
str(train)
```

```
## 'data.frame': 360336 obs. of 17 variables:
## $ id : int 0 1 2 3 4 5 6 7 8 9 ...
## $ store_sales.in.millions. : num 8.61 5 14.08 4.02 2.13 ...
## $ unit_sales.in.millions. : num 3 2 4 3 3 4 2 3 3 4 ...
## $ total_children : num 2 4 0 5 5 5 1 2 5 1 ...
## $ num_children_at_home : num 2 0 0 0 0 5 0 0 0 0 ...
## $ avg_cars_at_home.approx..1: num 2 3 3 0 3 3 2 2 2 3 ...
## $ gross_weight : num 10.3 6.66 21.3 14.8 17 7.26 9.58 16.9 13.8 15.7 ...
## $ recyclable_package : num 1 1 1 0 1 0 0 1 1 1 ...
## $ low_fat : num 0 0 0 1 1 1 0 0 0 1 ...
## $ units_per_case : num 32 1 26 36 20 5 6 2 6 9 ...
## $ store_sqft : num 36509 28206 21215 21215 27694 ...
## $ coffee_bar : num 0 1 1 1 1 1 1 1 0 1 ...
## $ video_store : num 0 0 0 0 1 0 1 1 0 1 ...
## $ salad_bar : num 0 0 0 0 1 1 1 1 0 1 ...
## $ prepared_food : num 0 0 0 0 1 1 1 1 0 1 ...
## $ florist : num 0 0 0 0 1 1 1 1 0 1 ...
## $ cost : num 62.1 121.8 83.5 66.8 111.5 ...
```

```
summary(test)
```

```
## id store_sales.in.millions. unit_sales.in.millions.
## Min. :360336 Min. : 0.510 Min. :1.000
## 1st Qu.:420392 1st Qu.: 3.750 1st Qu.:3.000
## Median :480448 Median : 5.800 Median :3.000
```

```
## Mean :480448 Mean : 6.354 Mean :3.044
## 3rd Qu.:540503 3rd Qu.: 8.400 3rd Qu.:4.000
## Max. :600559 Max. :22.920 Max. :6.000
## total_children num_children_at_home avg_cars_at.home.approx..1
## Min. :0.000 Min. :0.0000 Min. :0.000
## 1st Qu.:1.000 1st Qu.:0.0000 1st Qu.:1.000
## Median :2.000 Median :0.0000 Median :2.000
## Mean :2.454 Mean :0.6854 Mean :2.198
## 3rd Qu.:4.000 3rd Qu.:1.0000 3rd Qu.:3.000
## Max. :5.000 Max. :5.0000 Max. :4.000
## gross_weight recyclable_package low_fat units_per_case
## Min. : 6.00 Min. :0.0000 Min. :0.0000 Min. : 1.00
## 1st Qu.: 9.71 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:10.00
## Median :13.60 Median :1.0000 Median :0.0000 Median :20.00
## Mean :13.83 Mean :0.5657 Mean :0.3269 Mean :18.96
## 3rd Qu.:17.80 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:28.00
## Max. :21.90 Max. :1.0000 Max. :1.0000 Max. :36.00
## store_sqft coffee_bar video_store salad_bar
## Min. :20319 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:23593 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000
## Median :27694 Median :1.0000 Median :0.0000 Median :1.0000
## Mean :28175 Mean :0.5642 Mean :0.2756 Mean :0.5044
## 3rd Qu.:33858 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:1.0000
## Max. :39696 Max. :1.0000 Max. :1.0000 Max. :1.0000
## prepared_food florist
## Min. :0.0000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.0000
## Median :1.0000 Median :1.0000
## Mean :0.5044 Mean :0.5018
## 3rd Qu.:1.0000 3rd Qu.:1.0000
## Max. :1.0000 Max. :1.0000
```

```
str(test)
```

```
## 'data.frame': 240224 obs. of 16 variables:
## $ id : int 360336 360337 360338 360339 360340 360341 360342 360343 360344 360345 ...
## $ store_sales.in.millions. : num 7.24 6.9 8.34 5.48 4.8 5.25 3.72 7.68 9.63 9.44 ...
## $ unit_sales.in.millions. : num 4 2 3 2 3 3 4 4 3 4 ...
## $ total_children : num 1 2 0 3 2 1 4 2 3 5 ...
## $ num_children_at_home : num 0 2 0 3 0 1 0 0 1 0 ...
## $ avg_cars_at.home.approx..1: num 2 3 3 2 2 3 4 3 4 3 ...
## $ gross_weight : num 10.8 8.51 8.77 21.9 10.9 13.2 14.4 19.9 7.87 8.83 ...
## $ recyclable_package : num 0 1 0 1 1 1 1 0 1 1 ...
## $ low_fat : num 1 0 1 0 0 0 0 0 0 1 ...
## $ units_per_case : num 7 4 14 9 11 9 4 20 6 29 ...
## $ store_sqft : num 20319 33858 39696 23688 27694 ...
## $ coffee_bar : num 0 1 0 1 1 1 1 1 1 0 ...
## $ video_store : num 0 0 0 1 1 1 0 1 0 0 ...
## $ salad_bar : num 0 1 1 1 1 1 1 1 1 0 ...
## $ prepared_food : num 0 1 1 1 1 1 1 1 1 0 ...
## $ florist : num 0 1 0 1 1 1 1 1 1 1 ...
```

recyclable_package, low_fat, coffee_bar, video_store, salad_bar, prepared_food, and florist are binary variables.

id should be omitted from any model.

Look for correlation between variables.

```
cor(train[, -1])
```

```
##                                store_sales.in.millions. unit_sales.in.millions.
## store_sales.in.millions.                1.0000000000          0.4813757652
## unit_sales.in.millions.                0.481375765          1.0000000000
## total_children                        0.069303226          0.1132311894
## num_children_at_home                  0.029261049          0.0467545583
## avg_cars_at.home.approx..1            0.006794259          0.0171189249
## gross_weight                          0.038158576          0.0004667436
## recyclable_package                    0.030389878          0.0010739744
## low_fat                              -0.008735475          -0.0036626600
## units_per_case                       -0.009893103          0.0011405670
## store_sqft                           0.021571915          0.0455398313
## coffee_bar                           -0.040039499          -0.0769673037
## video_store                          0.029564152          0.0537948590
## salad_bar                            0.044831915          0.0824451025
## prepared_food                        0.044854158          0.0824847913
## florist                              0.046471926          0.0830621888
## cost                                -0.012386967          -0.0265087663
##                                total_children num_children_at_home
## store_sales.in.millions.      0.0693032262          0.0292610494
## unit_sales.in.millions.      0.1132311894          0.0467545583
## total_children                1.0000000000          0.3592070613
## num_children_at_home          0.3592070613          1.0000000000
## avg_cars_at.home.approx..1    0.0785191739          0.1154756738
## gross_weight                  -0.0009077197          -0.0006014972
## recyclable_package            0.0022356359          0.0061042640
## low_fat                      -0.0015974132          -0.0002076602
## units_per_case                -0.0002668987          -0.0041741829
## store_sqft                    -0.0089908507          0.0057847790
## coffee_bar                    -0.0064764368          -0.0232337747
## video_store                   -0.0133034887          -0.0207378304
## salad_bar                     -0.0235637940          -0.0311088116
## prepared_food                 -0.0235604814          -0.0310503584
## florist                       -0.0125453620          -0.0178879497
## cost                          -0.0074816998          -0.0017271315
##                                avg_cars_at.home.approx..1 gross_weight
## store_sales.in.millions.      0.006794259  0.0381585756
## unit_sales.in.millions.      0.017118925  0.0004667436
## total_children                0.078519174 -0.0009077197
## num_children_at_home          0.115475674 -0.0006014972
## avg_cars_at.home.approx..1    1.000000000 -0.0022671343
## gross_weight                  -0.002267134  1.0000000000
## recyclable_package            0.004020787  0.0590504265
## low_fat                      -0.001912331 -0.0334207253
## units_per_case                0.001190268 -0.0176090654
## store_sqft                    -0.020031687 -0.0004551021
## coffee_bar                    -0.000440755  0.0008351537
## video_store                   0.012702076 -0.0008447703
## salad_bar                     -0.013102435  0.0021163179
```

## prepared_food		-0.013035045	0.0021547017
## florist		-0.004284011	0.0007635955
## cost		0.027097743	-0.0001161770
##	recyclable_package	low_fat	units_per_case
## store_sales.in.millions.	0.0303898784	-0.0087354751	-0.0098931034
## unit_sales.in.millions.	0.0010739744	-0.0036626600	0.0011405670
## total_children	0.0022356359	-0.0015974132	-0.0002668987
## num_children_at_home	0.0061042640	-0.0002076602	-0.0041741829
## avg_cars_at.home.approx..1	0.0040207874	-0.0019123314	0.0011902682
## gross_weight	0.0590504265	-0.0334207253	-0.0176090654
## recyclable_package	1.0000000000	-0.0300252967	-0.0030280744
## low_fat	-0.0300252967	1.0000000000	0.0302257210
## units_per_case	-0.0030280744	0.0302257210	1.0000000000
## store_sqft	-0.0003428107	0.0019716790	0.0022966522
## coffee_bar	0.0040368730	0.0026217321	0.0008282796
## video_store	0.0041285191	0.0028927992	0.0005688030
## salad_bar	0.0046735853	0.0056020164	0.0016977467
## prepared_food	0.0046723238	0.0055671342	0.0016989680
## florist	0.0048485629	0.0055401218	0.0005564609
## cost	-0.0014548686	-0.0019750435	0.0001803538
##	store_sqft	coffee_bar	video_store
## store_sales.in.millions.	0.0215719146	-0.0400394985	0.0295641521
## unit_sales.in.millions.	0.0455398313	-0.0769673037	0.0537948590
## total_children	-0.0089908507	-0.0064764368	-0.0133034887
## num_children_at_home	0.0057847790	-0.0232337747	-0.0207378304
## avg_cars_at.home.approx..1	-0.0200316872	-0.0004407550	0.0127020757
## gross_weight	-0.0004551021	0.0008351537	-0.0008447703
## recyclable_package	-0.0003428107	0.0040368730	0.0041285191
## low_fat	0.0019716790	0.0026217321	0.0028927992
## units_per_case	0.0022966522	0.0008282796	0.0005688030
## store_sqft	1.0000000000	-0.1982428101	-0.0838731679
## coffee_bar	-0.1982428101	1.0000000000	0.5438258345
## video_store	-0.0838731679	0.5438258345	1.0000000000
## salad_bar	0.3330547104	0.4812480973	0.6136401003
## prepared_food	0.3331023373	0.4812250020	0.6136094442
## florist	-0.0741569709	0.5541826640	0.6154645610
## cost	-0.0492006363	-0.0520856748	-0.1067861697
##	salad_bar	prepared_food	florist
## store_sales.in.millions.	0.044831915	0.044854158	0.0464719263
## unit_sales.in.millions.	0.082445102	0.082484791	0.0830621888
## total_children	-0.023563794	-0.023560481	-0.0125453620
## num_children_at_home	-0.031108812	-0.031050358	-0.0178879497
## avg_cars_at.home.approx..1	-0.013102435	-0.013035045	-0.0042840111
## gross_weight	0.002116318	0.002154702	0.0007635955
## recyclable_package	0.004673585	0.004672324	0.0048485629
## low_fat	0.005602016	0.005567134	0.0055401218
## units_per_case	0.001697747	0.001698968	0.0005564609
## store_sqft	0.333054710	0.333102337	-0.0741569709
## coffee_bar	0.481248097	0.481225002	0.5541826640
## video_store	0.613640100	0.613609444	0.6154645610
## salad_bar	1.000000000	0.999839025	0.5986530671
## prepared_food	0.999839025	1.000000000	0.5986474858
## florist	0.598653067	0.598647486	1.0000000000
## cost	-0.098810123	-0.098843199	-0.1104140444

```
##                                cost
## store_sales.in.millions.    -0.0123869670
## unit_sales.in.millions.    -0.0265087663
## total_children              -0.0074816998
## num_children_at_home        -0.0017271315
## avg_cars_at.home.approx..1  0.0270977425
## gross_weight                -0.0001161770
## recyclable_package          -0.0014548686
## low_fat                     -0.0019750435
## units_per_case              0.0001803538
## store_sqft                  -0.0492006363
## coffee_bar                  -0.0520856748
## video_store                 -0.1067861697
## salad_bar                   -0.0988101229
## prepared_food               -0.0988431992
## florist                     -0.1104140444
## cost                        1.0000000000
```

salad_bar and prepared_food are highly correlated.

Initial linear regression.

```
lmmod = lm(cost ~ . - id, data = train)
summary(lmmod)
```

```
##
## Call:
## lm(formula = cost ~ . - id, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -62.445 -26.422   0.094  26.098  58.070
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      1.115e+02  4.196e-01  265.615 < 2e-16 ***
## store_sales.in.millions.  1.651e-02  1.706e-02   0.968  0.3333
## unit_sales.in.millions. -4.738e-01  7.311e-02 -6.481 9.11e-11 ***
## total_children        -1.925e-01  3.579e-02 -5.380 7.45e-08 ***
## num_children_at_home   -7.286e-02  4.378e-02 -1.664  0.0961 .
## avg_cars_at.home.approx..1  7.718e-01  4.592e-02  16.809 < 2e-16 ***
## gross_weight          -1.087e-03  1.074e-02 -0.101  0.9193
## recyclable_package     -6.174e-02  9.997e-02 -0.618  0.5369
## low_fat                -8.444e-02  1.054e-01 -0.801  0.4229
## units_per_case         1.159e-03  4.839e-03   0.240  0.8106
## store_sqft             -3.061e-04  1.028e-05 -29.764 < 2e-16 ***
## coffee_bar             1.045e+00  1.344e-01   7.776 7.52e-15 ***
## video_store            -5.151e+00  1.608e-01 -32.036 < 2e-16 ***
## salad_bar              6.628e+00  5.506e+00   1.204  0.2287
## prepared_food          -5.969e+00  5.506e+00 -1.084  0.2784
## florist                -4.958e+00  1.421e-01 -34.888 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 29.65 on 360320 degrees of freedom
## Multiple R-squared:  0.01953,    Adjusted R-squared:  0.01949
## F-statistic: 478.5 on 15 and 360320 DF,  p-value: < 2.2e-16
```

salad_bar has a higher absolute coefficient, so I will exclude prepared_food from future models.

Look for interactions.

```
lmmod2 = lm(cost ~ (unit_sales.in.millions. +
                    total_children + avg_cars_at.home.approx..1 +
                    store_sqft + coffee_bar + video_store +
                    florist)^2, data = train)
summary(lmmod2)
```

```
##
## Call:
## lm(formula = cost ~ (unit_sales.in.millions. + total_children +
##      avg_cars_at.home.approx..1 + store_sqft + coffee_bar + video_store +
##      florist)^2, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -61.148 -25.007   0.818  26.145  70.951
##
## Coefficients: (1 not defined because of singularities)
##                                     Estimate Std. Error
## (Intercept)                        1.073e+02  1.382e+00
## unit_sales.in.millions.            -8.868e-01  3.922e-01
## total_children                      3.089e-01  2.278e-01
## avg_cars_at.home.approx..1         -2.209e-01  3.031e-01
## store_sqft                         -1.154e-04  4.139e-05
## coffee_bar                         -8.066e+00  1.134e+00
## video_store                        -1.800e+01  8.012e+00
## florist                            2.242e+01  1.292e+00
## unit_sales.in.millions.:total_children -6.934e-02  4.280e-02
## unit_sales.in.millions.:avg_cars_at.home.approx..1  2.936e-02  6.024e-02
## unit_sales.in.millions.:store_sqft      1.566e-05  1.148e-05
## unit_sales.in.millions.:coffee_bar      -2.370e-01  1.587e-01
## unit_sales.in.millions.:video_store     -5.543e-01  2.005e-01
## unit_sales.in.millions.:florist         1.137e+00  1.721e-01
## total_children:avg_cars_at.home.approx..1  1.135e-01  3.086e-02
## total_children:store_sqft             -1.820e-05  5.705e-06
## total_children:coffee_bar             5.290e-02  8.828e-02
## total_children:video_store             1.114e+00  1.010e-01
## total_children:florist                -7.127e-01  9.186e-02
## avg_cars_at.home.approx..1:store_sqft    4.690e-06  7.850e-06
## avg_cars_at.home.approx..1:coffee_bar    7.849e-01  1.191e-01
## avg_cars_at.home.approx..1:video_store   -2.247e+00  1.356e-01
## avg_cars_at.home.approx..1:florist       1.275e+00  1.230e-01
## store_sqft:coffee_bar                 3.869e-04  3.853e-05
## store_sqft:video_store                 -2.711e-04  4.073e-05
## store_sqft:florist                    -1.316e-03  4.495e-05
## coffee_bar:video_store                  NA          NA
## coffee_bar:florist                     5.264e+00  3.438e-01
```

```
## video_store:florist                2.155e+01  7.898e+00
##                                     t value Pr(>|t|)
## (Intercept)                        77.645 < 2e-16 ***
## unit_sales.in.millions.            -2.261 0.023762 *
## total_children                      1.356 0.175236
## avg_cars_at.home.approx..1         -0.729 0.466143
## store_sqft                         -2.788 0.005309 **
## coffee_bar                         -7.114 1.13e-12 ***
## video_store                       -2.246 0.024682 *
## florist                           17.354 < 2e-16 ***
## unit_sales.in.millions.:total_children -1.620 0.105262
## unit_sales.in.millions.:avg_cars_at.home.approx..1 0.487 0.626003
## unit_sales.in.millions.:store_sqft 1.364 0.172450
## unit_sales.in.millions.:coffee_bar -1.494 0.135275
## unit_sales.in.millions.:video_store -2.765 0.005699 **
## unit_sales.in.millions.:florist 6.609 3.87e-11 ***
## total_children:avg_cars_at.home.approx..1 3.677 0.000236 ***
## total_children:store_sqft -3.190 0.001423 **
## total_children:coffee_bar 0.599 0.549064
## total_children:video_store 11.035 < 2e-16 ***
## total_children:florist -7.758 8.63e-15 ***
## avg_cars_at.home.approx..1:store_sqft 0.597 0.550214
## avg_cars_at.home.approx..1:coffee_bar 6.591 4.38e-11 ***
## avg_cars_at.home.approx..1:video_store -16.563 < 2e-16 ***
## avg_cars_at.home.approx..1:florist 10.369 < 2e-16 ***
## store_sqft:coffee_bar 10.042 < 2e-16 ***
## store_sqft:video_store -6.656 2.83e-11 ***
## store_sqft:florist -29.280 < 2e-16 ***
## coffee_bar:video_store NA NA
## coffee_bar:florist 15.313 < 2e-16 ***
## video_store:florist 2.728 0.006367 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 29.53 on 360308 degrees of freedom
## Multiple R-squared:  0.02751, Adjusted R-squared:  0.02743
## F-statistic: 377.4 on 27 and 360308 DF, p-value: < 2.2e-16
```

Make final linear regression model.

```
lmmod3 = lm(cost ~ unit_sales.in.millions. + total_children + avg_cars_at.home.approx..1 +
  store_sqft + coffee_bar + video_store +
  florist + unit_sales.in.millions.:video_store +
  unit_sales.in.millions.:florist + total_children:avg_cars_at.home.approx..1 +
  total_children:store_sqft + total_children:video_store +
  total_children:florist + avg_cars_at.home.approx..1:coffee_bar +
  avg_cars_at.home.approx..1:video_store +
  avg_cars_at.home.approx..1:florist + store_sqft:coffee_bar +
  store_sqft:video_store + store_sqft:florist + coffee_bar:florist +
  video_store:florist, data = train)
summary(lmmod3)
```

```
##
```



```
## Call:
## lm(formula = cost ~ unit_sales.in.millions. + total_children +
##     avg_cars_at.home.approx..1 + store_sqft + coffee_bar + video_store +
##     florist + unit_sales.in.millions.:video_store + unit_sales.in.millions.:florist +
##     total_children:avg_cars_at.home.approx..1 + total_children:store_sqft +
##     total_children:video_store + total_children:florist + avg_cars_at.home.approx..1:coffee_bar +
##     avg_cars_at.home.approx..1:video_store + avg_cars_at.home.approx..1:florist +
##     store_sqft:coffee_bar + store_sqft:video_store + store_sqft:florist +
##     coffee_bar:florist + video_store:florist, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -61.294 -25.034   0.813  26.135  71.230
##
## Coefficients:
##
##              Estimate Std. Error t value
## (Intercept)      1.063e+02  6.489e-01 163.808
## unit_sales.in.millions.      -6.655e-01  8.834e-02  -7.533
## total_children       1.134e-01  1.780e-01   0.637
## avg_cars_at.home.approx..1    3.228e-03  1.048e-01   0.031
## store_sqft        -5.664e-05  1.749e-05  -3.238
## coffee_bar       -8.724e+00  9.509e-01  -9.175
## video_store      -1.786e+01  8.012e+00  -2.229
## florist           2.273e+01  1.261e+00  18.022
## unit_sales.in.millions.:video_store    -6.492e-01  1.926e-01  -3.371
## unit_sales.in.millions.:florist       1.091e+00  1.677e-01   6.502
## total_children:avg_cars_at.home.approx..1  1.135e-01  3.063e-02   3.707
## total_children:store_sqft     -1.795e-05  5.519e-06  -3.253
## total_children:video_store    1.133e+00  9.588e-02  11.821
## total_children:florist      -7.072e-01  8.636e-02  -8.189
## avg_cars_at.home.approx..1:coffee_bar    7.654e-01  1.160e-01   6.597
## avg_cars_at.home.approx..1:video_store   -2.245e+00  1.354e-01 -16.584
## avg_cars_at.home.approx..1:florist     1.287e+00  1.221e-01  10.548
## store_sqft:coffee_bar    3.948e-04  3.798e-05  10.394
## store_sqft:video_store   -2.707e-04  4.072e-05  -6.647
## store_sqft:florist     -1.324e-03  4.450e-05 -29.760
## coffee_bar:florist       5.167e+00  3.403e-01  15.185
## video_store:florist      2.164e+01  7.898e+00   2.740
##
##              Pr(>|t|)
## (Intercept)      < 2e-16 ***
## unit_sales.in.millions.    4.96e-14 ***
## total_children      0.52417
## avg_cars_at.home.approx..1  0.97543
## store_sqft         0.00120 **
## coffee_bar         < 2e-16 ***
## video_store        0.02581 *
## florist            < 2e-16 ***
## unit_sales.in.millions.:video_store    0.00075 ***
## unit_sales.in.millions.:florist       7.93e-11 ***
## total_children:avg_cars_at.home.approx..1  0.00021 ***
## total_children:store_sqft      0.00114 **
## total_children:video_store      < 2e-16 ***
## total_children:florist        2.65e-16 ***
## avg_cars_at.home.approx..1:coffee_bar    4.20e-11 ***
```

```
## avg_cars_at.home.approx..1:video_store    < 2e-16 ***
## avg_cars_at.home.approx..1:florist        < 2e-16 ***
## store_sqft:coffee_bar                    < 2e-16 ***
## store_sqft:video_store                   3.00e-11 ***
## store_sqft:florist                       < 2e-16 ***
## coffee_bar:florist                       < 2e-16 ***
## video_store:florist                      0.00614 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 29.53 on 360314 degrees of freedom
## Multiple R-squared:  0.02748,    Adjusted R-squared:  0.02742
## F-statistic: 484.8 on 21 and 360314 DF,  p-value: < 2.2e-16
```

Predict on test data.

```
cost = predict(lmmod3, test)
head(cost)
```

```
##           1           2           3           4           5           6
## 102.45754  97.35886 102.05310  99.66160  93.91672  98.67901
```

```
summary(train$cost)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   50.79   70.32   98.81   99.61  126.62  149.75
```

```
lm.guess = cbind.data.frame(test$id, cost)
head(lm.guess)
```

```
##   test$id      cost
## 1 360336 102.45754
## 2 360337  97.35886
## 3 360338 102.05310
## 4 360339  99.66160
## 5 360340  93.91672
## 6 360341  98.67901
```

```
write.csv(lm.guess, 'lm.csv')
```

Linear regression model Kaggle score = .315. (Lower is better.)

Make spline model.

```
spline1 = lm(cost ~ ns(unit_sales.in.millions.) + ns(total_children) + ns(avg_cars_at.home.approx..1) +
              ns(store_sqft) + ns(coffee_bar) + ns(video_store) +
              ns(flорist), data = train)
summary(spline1)
```

```
##
## Call:
```

```
## lm(formula = cost ~ ns(unit_sales.in.millions.) + ns(total_children) +
##      ns(avg_cars_at.home.approx..1) + ns(store_sqft) + ns(coffee_bar) +
##      ns(video_store) + ns(florist), data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -56.63 -26.40   0.08  25.98  57.70
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      104.6394     0.2070  505.535 < 2e-16 ***
## ns(unit_sales.in.millions.)    -2.6496     0.4017   -6.596 4.23e-11 ***
## ns(total_children)             -1.3465     0.2089   -6.445 1.16e-10 ***
## ns(avg_cars_at.home.approx..1)   3.8019     0.2280   16.672 < 2e-16 ***
## ns(store_sqft)                 -6.8329     0.2043  -33.438 < 2e-16 ***
## ns(coffee_bar)                  1.4846     0.1615    9.191 < 2e-16 ***
## ns(video_store)                -6.1002     0.1839  -33.178 < 2e-16 ***
## ns(florist)                   -5.9443     0.1669  -35.618 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 29.65 on 360328 degrees of freedom
## Multiple R-squared:  0.01947,    Adjusted R-squared:  0.01945
## F-statistic: 1022 on 7 and 360328 DF,  p-value: < 2.2e-16
```

spline1 R-squared = .0194

Predict using spline model.

```
cost = predict(spline1, test)
spline.guess = cbind.data.frame(test$id, cost)
write.csv(spline.guess, 'spline.csv')
```

Kaggle score for spline model = .316

Make GAM model.

```
gam.mod = gam(cost ~ s(unit_sales.in.millions.) + s(total_children) + s(avg_cars_at.home.approx..1) +
               s(store_sqft) + coffee_bar + video_store +
               florist, data = train)
summary(gam.mod)
```

```
##
## Call: gam(formula = cost ~ s(unit_sales.in.millions.) + s(total_children) +
##      s(avg_cars_at.home.approx..1) + s(store_sqft) + coffee_bar +
##      video_store + florist, data = train)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -59.7851 -25.9954   0.7405  25.0822  60.1984
##
## (Dispersion Parameter for gaussian family taken to be 868.053)
##
##      Null Deviance: 322993390 on 360335 degrees of freedom
## Residual Deviance: 312773384 on 360316 degrees of freedom
```

```
## AIC: 3460736
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##
##              Df      Sum Sq Mean Sq  F value    Pr(>F)
## s(unit_sales.in.millions.)      1    302402   302402   348.3679 < 2e-16 ***
## s(total_children)                1      4009    4009    4.6181 0.03164 *
## s(avg_cars_at.home.approx..1)    1   257498   257498   296.6391 < 2e-16 ***
## s(store_sqft)                   1   688068   688068   792.6564 < 2e-16 ***
## coffee_bar                      1   400708   400708   461.6174 < 2e-16 ***
## video_store                     1   743678   743678   856.7198 < 2e-16 ***
## florist                         1  1031405  1031405  1188.1824 < 2e-16 ***
## Residuals                      360316 312773384      868
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##
##              Npar Df    Npar F      Pr(F)
## (Intercept)
## s(unit_sales.in.millions.)      3    15.78 2.939e-10 ***
## s(total_children)                3    71.21 < 2.2e-16 ***
## s(avg_cars_at.home.approx..1)    3    56.94 < 2.2e-16 ***
## s(store_sqft)                   3 1379.90 < 2.2e-16 ***
## coffee_bar
## video_store
## florist
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Predict using GAM model.

```
cost = predict(gam.mod, test)
gam.guess = cbind.data.frame(test$id, cost)
write.csv(gam.guess, 'gam.csv')
summary(cost)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    87.29  96.85   99.66   99.63 103.33  111.92
```

Kaggle score = .314

Make random forest model.

```
rf1 = ranger(cost ~ . - id, data = train, importance = 'permutation')
```

```
## Growing trees.. Progress: 51%. Estimated remaining time: 29 seconds.
## Computing permutation importance.. Progress: 49%. Estimated remaining time: 32 seconds.
```

```
summary(rf1)
```

```
##              Length Class      Mode
```

```
## predictions          360336 -none-      numeric
## num.trees             1 -none-      numeric
## num.independent.variables 1 -none-      numeric
## mtry                  1 -none-      numeric
## min.node.size         1 -none-      numeric
## variable.importance   15 -none-      numeric
## prediction.error       1 -none-      numeric
## forest                7 ranger.forest list
## splitrule             1 -none-      character
## treetype              1 -none-      character
## r.squared             1 -none-      numeric
## call                  4 -none-      call
## importance.mode       1 -none-      character
## num.samples           1 -none-      numeric
## replace               1 -none-      logical
```

```
rf1$variable.importance
```

```
##   store_sales.in.millions.  unit_sales.in.millions.
##           10.703273           13.061615
##           total_children    num_children_at_home
##           51.926407           35.171591
##   avg_cars_at.home.approx..1    gross_weight
##           50.896392           2.061862
##           recyclable_package    low_fat
##           1.064264           1.174548
##           units_per_case        store_sqft
##           1.907494           235.792467
##           coffee_bar           video_store
##           109.183614           75.263146
##           salad_bar           prepared_food
##           184.223360           166.930836
##           florist
##           243.906487
```

```
rf1$r.squared
```

```
## [1] 0.1059016
```

R-squared = .106

Make new RF model.

```
set.seed(21)
rf2 = ranger(cost ~ . -id -recyclable_package -units_per_case -gross_weight
             -low_fat, data = train)
```

```
## Growing trees.. Progress: 68%. Estimated remaining time: 14 seconds.
```

```
rf2$r.squared
```

```
## [1] 0.1136328
```

R-squared = .114

Predict using new RF model.

```
cost = predict(rf2, test)
rf.guess = cbind.data.frame(test$id, cost)
write.csv(rf.guess, 'rf.csv')
```

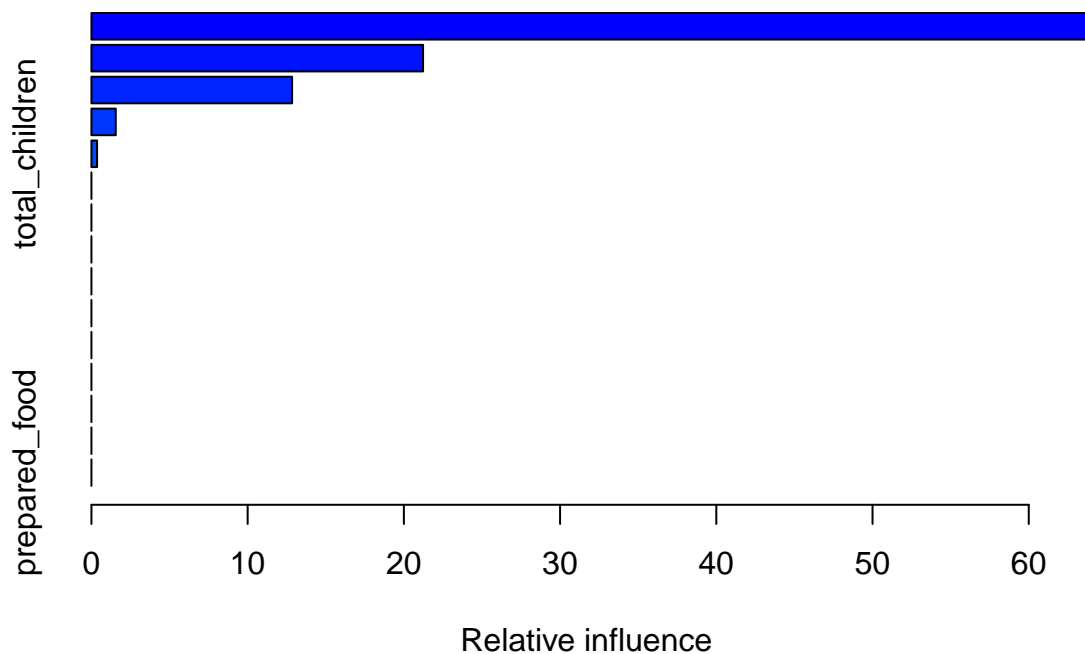
rf2 Kaggle score = .301

Make GBM models.

```
set.seed(21)
gbm1 = gbm(cost ~ . - id, data = train)
```

```
## Distribution not specified, assuming gaussian ...
```

```
summary(gbm1)
```



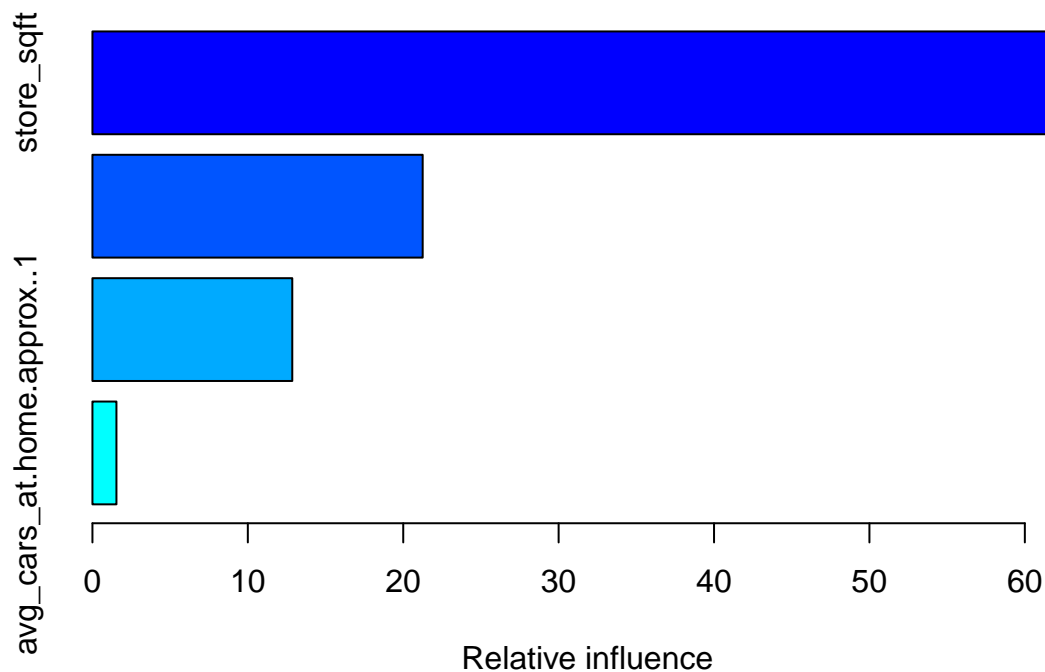
```
##               var    rel.inf
## store_sqft    store_sqft 64.0082788
## florist       florist  21.2292966
## video_store   video_store 12.8461256
## avg_cars_at.home.approx..1 avg_cars_at.home.approx..1 1.5563887
## total_children total_children 0.3599102
```

```
## store_sales.in.millions.    store_sales.in.millions.  0.0000000
## unit_sales.in.millions.    unit_sales.in.millions.  0.0000000
## num_children_at_home       num_children_at_home     0.0000000
## gross_weight               gross_weight     0.0000000
## recyclable_package         recyclable_package 0.0000000
## low_fat                   low_fat           0.0000000
## units_per_case             units_per_case    0.0000000
## coffee_bar                 coffee_bar       0.0000000
## salad_bar                  salad_bar        0.0000000
## prepared_food              prepared_food     0.0000000
```

```
set.seed(21)
gbm2 = gbm(cost ~ store_sqft + florist + video_store +
            avg_cars_at.home.approx..1, data = train)
```

```
## Distribution not specified, assuming gaussian ...
```

```
summary(gbm2)
```



```
##               var    rel.inf
## store_sqft      store_sqft 64.345276
## florist         florist  21.251004
## video_store     video_store 12.859261
## avg_cars_at.home.approx..1 avg_cars_at.home.approx..1  1.544458
```

Predict using GBM model.

```
cost = predict(gbm2, test)
```

```
## Using 100 trees...
```

```
gbm.guess = cbind.data.frame(test$id, cost)
write.csv(gbm.guess, 'gbm.csv')
```

Kaggle score for GBM model = .314

Make SVM model. Default kernel is radial.

```
train.svm = train[1:2402,]
```

```
svm1 = svm(cost ~ . - id, data = train.svm, kernel = 'linear')
```

Predict using SVM model.

```
cost = predict(svm1, test)
svm.guess = cbind.data.frame(test$id, cost)
write.csv(svm.guess, 'svm.csv')
```

SVM Kaggle score = .319

Optimize the random forest model, which has been the best model so far.

```
set.seed(23)
rf3 = ranger(cost ~ . -id -recyclable_package -units_per_case -gross_weight
             -low_fat, data = train, num.trees = 1000)
```

```
## Growing trees.. Progress: 34%. Estimated remaining time: 1 minute, 0 seconds.
## Growing trees.. Progress: 69%. Estimated remaining time: 28 seconds.
```

```
rf3$r.squared
```

```
## [1] 0.1132479
```

R-squared = .113. Adding more trees didn't help.

Try different values of mtry.

```
set.seed(23)
rf4 = ranger(cost ~ . -id -recyclable_package -units_per_case -gross_weight
             -low_fat, data = train, mtry = 1)
rf4$r.squared
```

```
## [1] 0.03702256
```

R-squared of rf4 = .038


```
set.seed(23)
rf5 = ranger(cost ~ . -id -recyclable_package -units_per_case -gross_weight
             -low_fat, data = train, mtry = 2)
rf5$r.squared
```

```
## [1] 0.0788579
```

```
rf5 R-squared = .079
```

```
set.seed(23)
rf6 = ranger(cost ~ . -id -recyclable_package -units_per_case -gross_weight
             -low_fat, data = train, mtry = 4)
```

```
## Growing trees.. Progress: 47%. Estimated remaining time: 35 seconds.
## Growing trees.. Progress: 95%. Estimated remaining time: 3 seconds.
```

```
rf6$r.squared
```

```
## [1] 0.1283611
```

```
rf6 R-squared = .128
```

```
set.seed(23)
rf7 = ranger(cost ~ . -id -recyclable_package -units_per_case -gross_weight
             -low_fat, data = train, mtry = 8)
```

```
## Growing trees.. Progress: 20%. Estimated remaining time: 2 minutes, 7 seconds.
## Growing trees.. Progress: 41%. Estimated remaining time: 1 minute, 30 seconds.
## Growing trees.. Progress: 62%. Estimated remaining time: 57 seconds.
## Growing trees.. Progress: 83%. Estimated remaining time: 25 seconds.
```

```
rf7$r.squared
```

```
## [1] 0.05681667
```

```
rf7 R-squared = .057
```

The best R-squared was achieved with an mtry of 4.

Use rf6 to predict.

```
cost = predict(rf6, test)
final.guess = cbind.data.frame(test$id, cost)
write.csv(final.guess, 'rf6.csv')
```

rf6 Kaggle score = .298. (Lower is better.) As of March 23rd, 2023, this puts me in 151st place out of 386.